

aq8tvdefd

May 7, 2025

# 1 Homework 3: Text Analysis of Bloomberg Articles

## 1.1 Data Cleaning and EDA

## 1.2 This Assignment

Welcome to Homework 3! For this assignment, we will work with Bloomberg news articles on Microsoft and Microsoft stock data (MSFT).

In this assignment, you will gain practice with:

- Conducting data cleaning and EDA on a text-based dataset,
- Manipulating data in `pandas` with the `datetime` and `string` accessors,
- Writing regular expressions and using `pandas` RegEx methods, and
- Performing sentiment analysis on text using DistilBERT.

```
[63]: # Run this cell to set up your notebook.
import warnings
warnings.simplefilter(action="ignore")

import re
import itertools
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from ds100_utils import *

# Ensure that pandas shows at least 280 characters in columns, so we can see
↳ full articles.
pd.set_option("max_colwidth", 280)
plt.style.use("fivethirtyeight")
sns.set()
sns.set_context("talk")
```

In this assignment, we will use the DistilBERT model, which is a Natural Language Processing (NLP) model designed to understand human language by processing text to capture the context and meaning of words within sentences. You are not expected to know the details of the model, but we will use it in this homework to perform sentiment analysis on textual data. We are importing

those tools and the corresponding model below. **If you see any warnings, please ignore them. As long as the cell runs, it shouldn't be any issues.**

```
[64]: from transformers import pipeline
      model_checkpoint = "distilbert/distilbert-base-uncased-finetuned-sst-2-english"
```

### 1.2.1 Score Breakdown

Question	Manual	Points
1a	No	1
1b	No	1
1c	No	3
1d	Yes	1
2a	No	2
2b	No	1
2c	No	2
2di	No	1
2dii	Yes	1
3ai	No	1
3aii	No	1
3b	No	2
3ci	No	1
3cii	Yes	1
<b>Total</b>	<b>3</b>	<b>19</b>

## 1.3 Question 1: Importing the Data

The data for this assignment is a subset of the financial news dataset from [this github repo](#). The original datasets are no longer available online due to copyright issues, but we were allowed access for educational purposes. The data in the file `data/msft_bloomberg_news.txt` has been filtered to just Bloomberg articles published between 2010 to 2013 (inclusive) with text that contains “Microsoft” or “MSFT” (Microsoft’s stock name).

---

### 1.3.1 Question 1a

Let’s examine the contents of the `data/msft_bloomberg_news.txt` file. Using the `open` function and `read` operation on a python file object, read **the first 1000 characters** in `data/msft_bloomberg_news.txt` and store your result in the variable `q1a`. Then, display the result so you can read it.

**CAUTION: Viewing the contents of large files in a Jupyter Notebook could crash your browser. Be careful not to print the entire contents of the file.**

```
[65]: q1a = open('/content/msft_bloomberg_news.txt', 'r')

      q1a = q1a.read(1000)
```

```
print(q1a)
```

```
[{"id":46243185,"title":"Opera Jumps Most Ever After Report Facebook May Bid:
Oslo Mover","released_at":"<date>May 29 2012</date>
<time>09:40:58</time>","content":"Opera Software ASA (OPERA) , the
Norwegian\marker of Internet browsers, surged the most on record in Oslo\after
technology website Pocket-Lint reported that Facebook Inc. (FB) \nmay try to
acquire the company. Opera gained as much as 26 percent, the biggest jump
since\nit first sold shares in 2004. The Oslo-based company rose 18\npercent to
40.5 kroner at 11:37 a.m., giving it a market value\nof 4.85 billion kroner
($807 million). Opera is the last major independent browser left, with
the\nothers owned by companies such as Microsoft Corp. (MSFT) , Google Inc.
(GOOG) \nand Apple Inc. (AAPL) , said Aleksander Nilsen, an analyst at Abg
Sundal\nCollier in Oslo. The company has a strong balance sheet, and\ncould be
an attractive target for other companies, such as\nMountain View , California-
based Google, he said.
```

---

### 1.3.2 Question 1b

Based on the printed output you got from q1a, what format is the data in? Answer this question by entering the letter corresponding to the right format in the variable q1b below.

**CAUTION:** As a reminder, viewing the contents of large files in a Jupyter Notebook could crash your browser. Be careful not to print the entire contents of the file, and do not use the file explorer to open data files directly.

A. CSV B. HTML C. JavaScript Object Notation (JSON) D. Excel XML

Answer in the following cell. Your answer should be a string, either "A", "B", "C", or "D".

```
[66]: q1b = 'C'
```

---

### 1.3.3 Question 1c

pandas has built-in readers for many different file formats, including the file format used here to store news articles. To learn more about these, check out the documentation for

- `pd.read_csv` ([docs](#))
- `pd.read_html` ([docs](#))
- `pd.read_json` ([docs](#))
- `pd.read_excel` ([docs](#)).

For this question, use one of these functions to: 1. Load the file `msft_bloomberg_news.txt` in the data folder as a `DataFrame` into the variable `msft_news_df`. 2. Set the `index` of `msft_news_df` to correspond to the `id` of each news article.

**Hint:** If your code is taking a while to run, you should review your answers to q1a and q1b; you may have used the incorrect data loading function for the type of the given file.

```
[67]: msft_news_df = pd.read_json('/content/msft_bloomberg_news.txt')
      msft_news_df.head(1)
```

```
[67]:      id                                     title \
0  46243185  Opera Jumps Most Ever After Report Facebook May Bid: Oslo Mover

      released_at \
0  <date>May 29 2012</date> <time>09:40:58</time>

      content \
0  Opera Software ASA (OPERA) , the Norwegian\marker of Internet browsers,
surged the most on record in Oslo\after technology website Pocket-Lint
reported that Facebook Inc. (FB) \may try to acquire the company. Opera
gained as much as 26 percent, the biggest jump since\n...

      path
0  ./2008_2012_msft_bloomberg_news/opera-jumps-most-on-record-after-report-of-
facebook-s-interes.txt
```

---

### 1.3.4 Question 1d

Suppose we are interested in using the news to predict future stock values. What additional data would we need to predict stock prices, and how could we connect that data to news articles? In addition, what attributes or characteristics of the news might help predict the stock value?

Com base na questão 1, um dado adicional que poderia ajudar seria todo o histórico de mercado das empresas. E como o artigo fornece o horário em que a Opera aumentou o seu valor de mercado, o atributo hora pode ser útil para prever futuros valores de mercado, pois ações e valores de mercado estão sempre mudando

## 1.4 Question 2: Time Analysis

After loading in the data, we can start exploring news articles by analyzing the relationships between the release dates (date of publication) and different topics and companies.

---

### 1.4.1 Question 2a

First, let's extract the date and time from the `released_at` column in `msft_news_df`. Notice that the date and time are encoded in the following format:

```
<date>May 29 2012</date> <time>09:40:58</time>
<date>May 18 2011</date> <time>22:42:40</time>
<date>August 15 2012</date> <time>00:09:02</time>
<date>July 1 2011</date> <time>22:12:37</time>
...
```

There are several ways to convert this to a `Timestamp` object that we can use more easily. However, for this assignment, we are going to use string manipulation functions.

Create a regular expression that extracts the Month, Day, Year, Hour, Minute, and Second from the `msft_news_df["released_at"]` column. You should create a new `DataFrame` called `dates` that contains: 1. The same index as `msft_news_df` (`id`) and 2. Column labels: "Month", "Day", "Year", "Hour", "Minute", "Second".

Additionally, convert all numerical values ("Year", "Day", "Hour", "Minute", "Second") to type `int`.

**Hint 1:** You should use the `Series.str.extract` function.

**Hint 2:** Don't forget to use raw strings and capture groups. Copy the above example text into [regex101.com](http://regex101.com) to experiment with your regular expressions.

**Hint 3:** It might be helpful to break this up into a couple of steps (e.g., first extract date values such as Month, Day, and Year and then extract time values such as Hour, Minute, and Second).

```
[68]: dates = msft_news_df['released_at']

mes = dates.str.extract(r'<date>(P<letter>\w+)' )
dia = dates.str.extract(r'(\d+)').astype(int)
ano = dates.str.extract(r'(\d{4})</date>').astype(int)
hora = dates.str.extract(r'<time>(\d{2})').astype(int)
min = dates.str.extract(r':(\d{2}):').astype(int)
seg = dates.str.extract(r':(\d{2})</time>').astype(int)

dates = pd.DataFrame(dates)
dates['Month'] = mes
dates['Day'] = dia
dates['Year'] = ano
dates['Hour'] = hora
dates['Minute'] = min
dates['Second'] = seg

dates
```

```
[68]:
```

	released_at	Month	Day \
0	<date>May 29 2012</date> <time>09:40:58</time>	May	29
1	<date>May 18 2011</date> <time>22:42:40</time>	May	18
2	<date>August 15 2012</date> <time>00:09:02</time>	August	15
3	<date>July 1 2011</date> <time>22:12:37</time>	July	1
4	<date>January 18 2012</date> <time>01:20:28</time>	January	18
...	...	...	...
4630	<date>June 27 2012</date> <time>00:35:58</time>	June	27
4631	<date>September 24 2013</date> <time>13:38:57</time>	September	24
4632	<date>September 14 2011</date> <time>04:01:00</time>	September	14
4633	<date>June 28 2010</date> <time>01:00:00</time>	June	28
4634	<date>September 8 2011</date> <time>01:11:01</time>	September	8

	Year	Hour	Minute	Second
0	2012	9	40	58
1	2011	22	42	40
2	2012	0	9	2
3	2011	22	12	37
4	2012	1	20	28
...	...	...	...	...
4630	2012	0	35	58
4631	2013	13	38	57
4632	2011	4	1	0
4633	2010	1	0	0
4634	2011	1	11	1

[4635 rows x 5 columns]

### 1.4.2 Question 2b

Now that we've figured out how to extract dates, create a new `DataFrame` called `msft_news_2010` that only contains articles released in 2010. This `DataFrame` should contain: 1. An index of `id` and 2. Columns: "title", "released\_at", "content", "path", "Month", "Day", and "Year".

**Hint:** Consider merging `msft_news_df` with `dates`.

```
[69]: msft_news_2010 = msft_news_df.loc[dates['Year'] == 2010].merge(dates)
```

```
msft_news_2010
```

```
[69]:      id \
0    95357231
1    75227517
2    57850804
3    75532360
4    10176588
..      ...
573  95653167
574  44065090
575  12166320
576   9764270
577  25935811
```

```

title
\
0    Netflix Profit Jumps 44% on New Users
1    Republican Win May Be Tax Boon for Companies, High Incomes
2    Alibaba Says It Now Offers Sohu's Search Engine
```

3 Slim Solution for Trade Imbalances Is More Buying by China  
4 S&P 500 to Defy 'New Normal' and Rally 17%, Cambiar's Barish Says  
..  
573 Apple to Open Digital Store for Mac Computer Apps  
574 Buffett Donates \$1.6 Billion in Biggest Gift Since 2008 Crisis  
575 Nintendo Bars Children Under 6 From Viewing 3-D Images on New Game Player  
576 Microsoft's Ballmer Says Tablet Computers 'Top of Mind' Amid Apple Success  
577 Stocks With High Profit, Low Debt Keep Me Calm: John Dorfman

released\_at \

0 <date>April 21 2010</date> <time>23:52:36</time>  
1 <date>November 3 2010</date> <time>16:46:00</time>  
2 <date>October 29 2010</date> <time>12:23:43</time>  
3 <date>October 31 2010</date> <time>16:05:40</time>  
4 <date>December 1 2010</date> <time>20:38:58</time>  
..  
573 <date>December 16 2010</date> <time>21:57:29</time>  
574 <date>July 6 2010</date> <time>04:00:03</time>  
575 <date>December 30 2010</date> <time>01:47:51</time>  
576 <date>July 30 2010</date> <time>00:01:56</time>  
577 <date>June 28 2010</date> <time>01:00:00</time>

content \

0 Netflix Inc. said first-quarter\nprofit rose 44 percent as the movie  
subscription service signed\nup new customers and increased online offerings.  
\n Net income advanced to \$32.3 million, or 59 cents a share,\nfrom \$22.4  
million, or 37 cents, a year earlier, the Los Gatos,\n...  
1 Americans with the highest incomes\nand U.S. corporations, especially those  
with international\noperations, stand to be big winners as newly  
elected\ncongressional Republicans signal they will extend existing  
tax\nbenefits and push for new ones. Republicans will use their ne...  
2 Alibaba Group Holding Ltd. said\nusers of its search-engine service may  
now access technology\nsupplied by Sohu.com Inc. , as the two Chinese companies  
\nstrengthen collaboration to challenge industry leader Baidu Inc. Users of  
Alibaba's Etao.com search service may now o...  
3 Billionaire Carlos Slim , the world's\nrichest man, said China must buy  
more and the U.S. needs to step\nup private investment to reduce the trade  
imbalance and boost\ntheir economies. Global currency devaluation efforts will  
fail in the\nabsence of economic policies that f...  
4 Energy and industrial companies will\nrise next year, propelling a 17  
percent gain in the Standard &\nPoor's 500 Index from its current level,  
according to Cambiar\nInvestors LLC's Brian Barish . Next year will be marked  
by a "multi-speed recovery" as\nindustries weakened b...  
..  
...  
573 Apple Inc. will open a digital\nstorefront next month that will try to do  
for computer software\nwhat it did for music and mobile applications. The Mac

App Store will open Jan. 6, the Cupertino,\nCalifornia-based company said in a statement today. The aim is\nto let Mac own...

574 Warren Buffett , the billionaire who\nhas promised to give away 99 percent of his fortune to charity,\nmade his largest donation since the 2008 financial crisis after\nprofits at his Berkshire Hathaway Inc. jumped. \n The value of Buffett's annual gift to the foundation\ne...

575 Nintendo Co. will bar children ages\n6 and younger from using the 3-D functions of its new handheld\ngame machine at an introductory event for the device. "Looking at 3-D images for a long time may harm the growth\nof children's eyes," Nintendo said in a note to visitors ...

576 Microsoft Corp. Chief Executive\nOfficer Steve Ballmer said tablet computers are high on his\npriority list as Apple Inc. takes the lead in a market his\ncompany has tried to foster for more than a decade. \n "Today, one of the top issues on my mind is 'hey there's a\ncat...

577 With many investors nervous as March\nhares, this is a good time to look at stocks with high profits\nand low debt. \n I don't expect a double-dip recession, but if one happens,\nthese companies should withstand it better than most. They\nshould also do fine if the economy i...

path \

0 ./2008\_2012\_msft\_bloomberg\_news/netflix-quarterly-profit-increases-44-as-movie-rental-service-adds-users.txt

1 ./2008\_2012\_msft\_bloomberg\_news/republican-sweep-may-mean-tax-boon-for-u-s-multinationals-high-incomes.txt

2 ./2008\_2012\_msft\_bloomberg\_news/alibaba-says-it-now-offers-sohu-s-search-engine.txt

3 ./2008\_2012\_msft\_bloomberg\_news/slim-solution-for-trade-imbalances-is-more-buying-by-china-u-s-investing.txt

4 ./2008\_2012\_msft\_bloomberg\_news/s-p-500-to-defy-pimco-s-new-normal-rise-17-by-end-of-2011-barish-says.txt

..

...

573 ./2008\_2012\_msft\_bloomberg\_news/apple-aims-to-do-for-computer-software-what-it-did-for-mobile-music-apps.txt

574 ./2008\_2012\_msft\_bloomberg\_news/buffett-donates-most-since-2008-after-urging-wealthiest-to-increase-giving.txt

575 ./2008\_2012\_msft\_bloomberg\_news/nintendo-bars-children-under-6-from-viewing-3-d-images-on-new-game-player.txt

576 ./2008\_2012\_msft\_bloomberg\_news/icrosoft-s-ballmer-says-tablet-computers-top-of-mind-amid-apple-success.txt

577 ./2008\_2012\_msft\_bloomberg\_news/stocks-with-high-profit-low-debt-keep-me-calm-john-dorfman.txt

	Month	Day	Year	Hour	Minute	Second
0	April	21	2010	23	52	36
1	November	3	2010	16	46	0
2	October	29	2010	12	23	43



3	October	31	2010	16	5	40
4	December	1	2010	20	38	58
..	...	...	...	...	...	...
573	December	16	2010	21	57	29
574	July	6	2010	4	0	3
575	December	30	2010	1	47	51
576	July	30	2010	0	1	56
577	June	28	2010	1	0	0

[578 rows x 11 columns]

### 1.4.3 Question 2c

After processing the article release dates, we can analyze articles about different topics and companies. Note that all the articles in the provided dataset mention Microsoft/MSFT, but they can also mention other companies.

For each company in the list of `companies` (provided below), add a boolean column to the `msft_news_df` DataFrame indicating whether the corresponding company is mentioned in the text of the article. Ultimately, you should add six new columns containing `True/False` values to the DataFrame: "amazon", "nintendo", "apple", "sony", "facebook", "netflix". You may use a `for` loop over the list of companies.

**Note:** Make the contents of the articles lowercase before searching for the keywords.

```
[70]: companies = ["amazon", "nintendo", "apple", "sony", "facebook", "netflix"]

msft_news_df['content'] = msft_news_df['content'].str.lower()

for i in companies:
    msft_news_df[i] = msft_news_df['content'].str.contains(i)

msft_news_df
```

```
[70]:      id \
0      46243185
1      73522879
2      29296500
3      49799724
4      20739032
...
4630   75325873
4631   49071474
4632   12417018
4633   25935811
4634   21143940
```

	title
\	
0	Opera Jumps Most Ever After Report Facebook May Bid: Oslo Mover
1	Microsoft Calls Intel's Comments on Next Windows 'Inaccurate'
2	Lawyers Raking in Cash as Campaign Spending Hits Records
3	Microsoft, Google Sued by Louisiana Firm Over Computer-Mapping Technology
4	Yahoo Co-Founder Jerry Yang Exits Company
...	...
4630	Dolby to Purchase San Francisco Tower for \$109.8 Million
4631	Mayfair Office Squeeze Spawns New London Real Estate Hubs
4632	Only Half of U.S. Corporate Cash Stays at Home: Chart of the Day
4633	Stocks With High Profit, Low Debt Keep Me Calm: John Dorfman
4634	Yahoo Shares Surge After Chairman Ends Bartz's Tenure With Telephone Call

	released_at	\
0	<date>May 29 2012</date>	<time>09:40:58</time>
1	<date>May 18 2011</date>	<time>22:42:40</time>
2	<date>August 15 2012</date>	<time>00:09:02</time>
3	<date>July 1 2011</date>	<time>22:12:37</time>
4	<date>January 18 2012</date>	<time>01:20:28</time>
...	...	
4630	<date>June 27 2012</date>	<time>00:35:58</time>
4631	<date>September 24 2013</date>	<time>13:38:57</time>
4632	<date>September 14 2011</date>	<time>04:01:00</time>
4633	<date>June 28 2010</date>	<time>01:00:00</time>
4634	<date>September 8 2011</date>	<time>01:11:01</time>

	content	\
0	opera software asa (opera) , the norwegian\marker of internet browsers, surged the most on record in oslo\after technology website pocket-lint reported that facebook inc. (fb) \may try to acquire the company. opera gained as much as 26 percent, the biggest jump since\...	
1	microsoft corp. (msft) said comments made by\an intel corp. (intc) executive yesterday about future version of its\windows operating system were "factually inaccurate and\nunfortunately misleading." renee james, head of intel's software business, said\yesterday that mi...	
2	every four years, a new mix of politicians assembles to compete for the opportunity to run for president. while the candidates' names and faces change, the lawyers stay the same. attorney michael toner began his presidential-campaign legal career in 1996 working for republic...	
3	microsoft corp. (msft) and google inc. (goog) were\naccused of violating a louisiana company's patent covering\nmapping technology that helps computer users see locations in\nthree dimensions. officials of transcenic inc. contend in a lawsuit that\nexecutives of google,...	
4	jerry yang is exiting the yahoo!\ninc (yhoo) . board and its management team, the latest casualty of an\noverhaul that led to the ouster of chief executive officer carol\nbartz and left the company in search of strategic	

options. yang, who started yahoo in 1995 with dav...

...

...

4630 dolby laboratories inc. (dlb) , the audio-\ntechnology company whose products are used in cinemas, recording\nstudios and video games, agreed to buy a 16-story tower in the\n san francisco area that's home to twitter inc., and will make\nthe building its new headquarters. t...

4631 mayfair and st. james's just aren't\nbig enough for all the companies that want a piece of london's\nmost expensive neighborhoods. many are now settling for less\nprestigious city-center addresses, creating new hot spots in the\noffice-property market. buildings are sproutin...

4632 cash levels for u.s. companies are\nlosing their meaning for the country's economy because so much\nof the money is held elsewhere these days, according to dane mott, a jpmorgan chase & co. analyst. mott drew his conclusion from a review of 258 companies\nthat disclose cash ...

4633 with many investors nervous as march\nhares, this is a good time to look at stocks with high profits\nand low debt. \n i don't expect a double-dip recession, but if one happens,\nthese companies should withstand it better than most. they\nshould also do fine if the economy i...

4634 yahoo! inc. surged after firing\nchief executive officer carol bartz, whose reign was marked by\nfalling sales, lost share to rivals and a dispute with asian\npartners that stunted growth in the world's largest web market. bartz said in a memo to staff yesterday that she was...

path \

0 ./2008\_2012\_msft\_bloomberg\_news/opera-jumps-most-on-record-after-report-of-facebook-s-interes.txt

1 ./2008\_2012\_msft\_bloomberg\_news/icrosoft-calls-intel-s-comments-on-next-windows-inaccurate-.txt

2 ./2008\_2012\_msft\_bloomberg\_news/awyers-raking-in-cash-as-campaign-spending-hits-records.txt

3 ./2008\_2012\_msft\_bloomberg\_news/icrosoft-google-sued-over-technology-providing-computer-maps.txt

4 ./2008\_2012\_msft\_bloomberg\_news/yahoo-says-co-founder-jerry-yang-resigns.txt

...

...

4630 ./2008\_2012\_msft\_bloomberg\_news/dolby-to-purchase-san-francisco-tower-for-109-8-million.txt

4631 ./2008\_2012\_msft\_bloomberg\_news/ayfair-office-squeeze-spawns-new-london-real-estate-hubs.txt

4632 ./2008\_2012\_msft\_bloomberg\_news/only-half-of-u-s-corporate-cash-stays-at-home-chart-of-the-day.txt

4633 ./2008\_2012\_msft\_bloomberg\_news/stocks-with-high-profit-low-debt-keep-me-calm-john-dorfman.txt

4634 ./2008\_2012\_msft\_bloomberg\_news/yahoo-s-carol-bartz-is-said-to-be-stepping-down-as-chief-executive-officer.txt

	amazon	nintendo	apple	sony	facebook	netflix
0	False	False	True	False	True	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	True	False
...	...	...	...	...	...	...
4630	False	False	False	False	False	False
4631	True	False	False	False	False	False
4632	False	False	True	False	False	False
4633	False	False	True	False	False	False
4634	False	False	False	False	True	False

[4635 rows x 11 columns]

#### 1.4.4 Question 2d

Now, we can put everything together to analyze the release dates and volume of articles for different companies.

**Question 2d, Part i** Create a new `DataFrame` called `year_news` that contains the number of articles mentioning each company in the list `companies` after 2010 (inclusive). `year_news` should have six columns (one column for each company), and the index of this `DataFrame` should be the release year "Year".

```
[71]: year_news = msft_news_df.loc[dates['Year'] >= 2010].merge(dates)
      year_news = year_news.groupby('Year')[companies].sum()

      year_news.head()
```

```
[71]:
```

	amazon	nintendo	apple	sony	facebook	netflix
Year						
2010	41	28	198	55	74	9
2011	102	29	491	105	163	43
2012	186	47	796	103	281	41
2013	158	94	723	201	254	51

**Question 2d, Part ii** Given your code in the previous part is correct, after running the cell below, you should be able to see the number of articles released mentioning `companies` for each year. The plot should look like this:

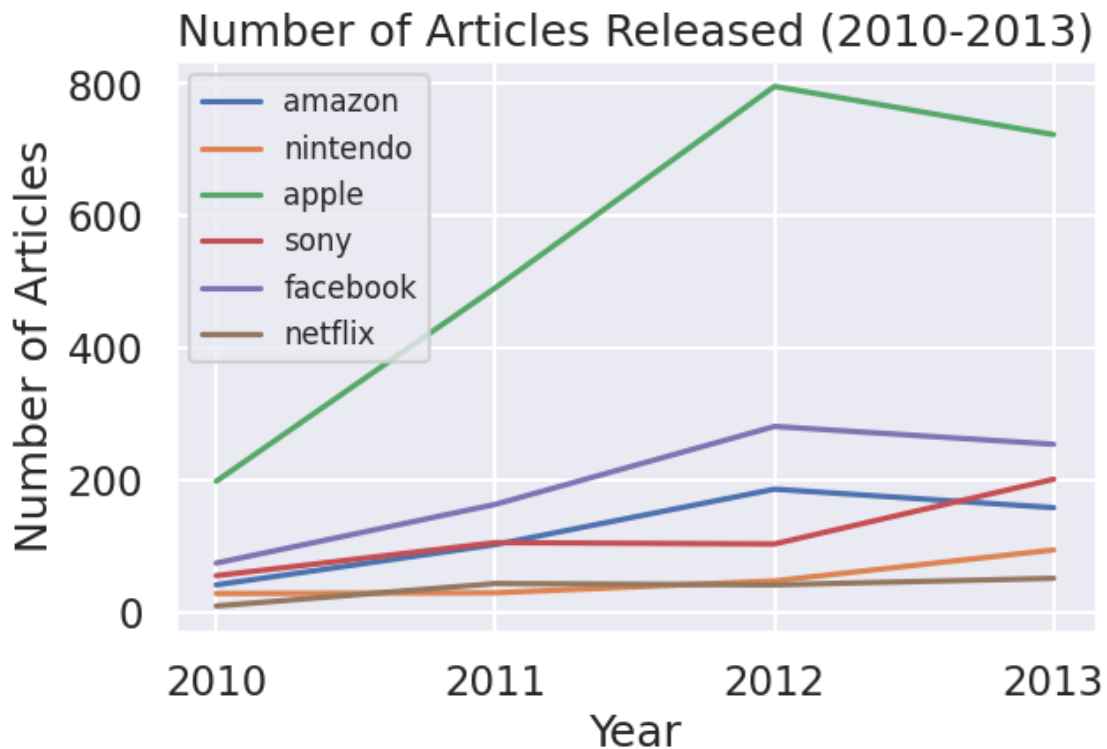
```
[72]: plt.figure(figsize=(6, 4))

      for company in companies:
```

```

sns.lineplot(data=year_news.reset_index(),
             x="Year",
             y=company,
             label=company)
plt.legend(fontsize="12")
plt.xticks(np.arange(2010, 2014), np.arange(2010, 2014))
plt.ylabel("Number of Articles")
plt.xlabel("Year")
plt.title("Number of Articles Released (2010-2013)");

```



What trends do you notice in the plot above? Feel free to reference or Google any events to explain the trends seen in the graph. What are some limitations of using data and the corresponding plot to analyze the performance of different companies or trends?

**Hint:** Remember the source of the articles and the subset of the articles we are analyzing in this assignment.

O gráfico demonstra o crescimento que as empresas tiveram depois de 2010, como pode ser percebido pelos dados da Apple que tiveram um pico mais alto em 2012 com o lançamento do iPhone 5, que foi o primeiro lançado depois da morte do Steve Jobs e que foi um sucesso de vendas

## 1.5 Question 3: Sentiment Analysis

In this section, we will continue building on our past analysis and specifically look at the **sentiment of each article** — this will lead us to a much more direct and detailed understanding of how these articles can be used in different applications. **Sentiment analysis** is generally the computational task of classifying the emotions in a body of text as positively or negatively charged.

We will use a fine-tuned version of the **DistilBERT** model ([github](#), [original paper](#)) to analyze the sentiment of Bloomberg news articles. DistilBERT is a neural network-based language model (a close relative to ChatGPT); we will use the model checkpoint specifically trained for sentiment analysis. These models are not in scope for Data 100, and we don't expect you to know how they work; take CS182: Neural Networks or Data 102: Data, Inference, and Decisions if you're interested in learning more. We are using them here to show how easy (and useful) these technologies have become.

We can use the [HuggingFace library](#) to build the sentiment analysis pipeline and load the model. [Here](#) is the card of the model checkpoint we will use for this assignment: the model card contains general information about the model, including the base model used, training arguments, training data, etc. Again, you don't need to know this for the course but knowing about model cards is important when you start to use these techniques in your careers.

Run the following two cells to set up the sentiment analysis pipeline and see examples of how we can get the sentiment for different strings.

```
[73]: # Load the model
sentiment_analysis = pipeline("sentiment-analysis", model=model_checkpoint)

# Get the sentiment of a given string
sentiment_1 = sentiment_analysis("I have two dogs.")
print("Example 1: " + str(sentiment_1))

sentiment_2 = sentiment_analysis("I do not have dogs.")
print("Example 2: " + str(sentiment_2))

sentiment_3 = sentiment_analysis("Fortunately, I do not have dogs to worry_
↪about.")
print("Example 3: " + str(sentiment_3))
```

Device set to use cpu

Example 1: [{'label': 'POSITIVE', 'score': 0.9955033659934998}]

Example 2: [{'label': 'NEGATIVE', 'score': 0.9987561702728271}]

Example 3: [{'label': 'POSITIVE', 'score': 0.9975079298019409}]

As you can see, the model can determine the sentiment of phrases/sentences (not just words). The model measures the phrase's **polarity**, indicating how strongly negative or positive it is on a scale of 0 to 1.

**Note:** The output is a list, and each element of the list is a dictionary with two keys (label and score). Note that we could have gotten the sentiments of the two sentences by putting them in a list (batch) and then running the pipeline once (see the code below).

```
[74]: sentiments = sentiment_analysis(["I have two dogs.", "I do not have dogs."])
      print(sentiments)
```

```
[{'label': 'POSITIVE', 'score': 0.9955033659934998}, {'label': 'NEGATIVE',
'score': 0.9987561702728271}]
```

---

### 1.5.1 Question 3a

As running all the articles through the model will take a while, let's first focus on articles released in 2010. We have already filtered these articles in `q2b` and assigned them to the `DataFrame` `msft_news_2010`.

Due to model input size constraints, a maximum of 512 words (tokens), and limited computational resources on Datahub, we cannot load the full articles into the pipeline. Instead, we can look at the first sentence that mentions Microsoft in each article.

**Question 3a, Part i** Assign `microsoft_re` to a regular expression that captures sentences referencing “microsoft” or “msft” (in lowercase). You should assume all sentences end with `.`, `?`, or `!` and that these punctuation characters are not used for any other purpose. This is of course not true in practice (e.g., this example! and 3.14), but we will often make these simplifying assumptions to enable progress in data analysis.

You should develop and test your regular expression using [regex101.com](https://regex101.com). Here are some practice sentences.

```
have you ever worked at microsoft? i once did. microsoft is known for
their research in ai.
```

Then: 1. Canonicalize the `"content"` of the articles by converting the text to lowercase, 2. Use the `microsoft_re` regular expression to extract the first sentence mentioning “microsoft” or “msft” in each article, and 3. Create a new column `first_sentence` in `msft_news_2010` with these values.

**Hint 1:** `Series.str.findall` function might be useful (might take around a minute to run).

**Hint 2:** Consider using the negation character class `r"[^.!?]"`

**Hint 3:** Some sentences will wrap across lines and the `.` will not match across new lines.

```
[75]: msft_news_2010['content'] = msft_news_2010['content'].str.lower()

microsoft_re = msft_news_2010['content'].str.extract(r'([^.?!]*\b(?:
↪microsoft|msft)\b([^.?!]*[.?!]))')[0]
microsoft_re = pd.DataFrame(microsoft_re)

msft_news_2010['first_sentence'] = microsoft_re
msft_news_2010.head(1)
```

```
[75]:      id                                     title \
0  95357231  Netflix Profit Jumps 44% on New Users
```

```

                                released_at \
0 <date>April 21 2010</date> <time>23:52:36</time>

                                content \
0 netflix inc. said first-quarter\nprofit rose 44 percent as the movie
subscription service signed\nup new customers and increased online offerings.
\n net income advanced to $32.3 million, or 59 cents a share,\nfrom $22.4
million, or 37 cents, a year earlier, the los gatos,\nn...

                                path \
0 ./2008_2012_msft_bloomberg_news/netflix-quarterly-profit-increases-44-as-
movie-rental-service-adds-users.txt

    Month Day Year Hour Minute Second \
0 April 21 2010 23 52 36

                                first_sentence
0 \n "if we had offered a pay-per-view service for new\nreleases, we would be
in conflict with a broad range of\ncompanies, including wal-mart, microsoft,
sony and apple,"nhastings said.

```

**Question 3a, Part ii** Using the `sentiment_analysis` model, let's now determine the sentiment of the first sentence that mentions "microsoft" or "msft" for each article. Note that the model outputs both a label and a score. Provide just the score, which should be converted to a negative number if the label is "NEGATIVE". Add a new column `sentence_sentiment` to `msft_news_2010` with these values.

**Note 1:** Feel free to reference the start of q3 to understand what `sentiment_analysis` can take in and what it outputs. `sentiment_analysis` may take 1-2 minutes to run when calculating scores for all the sentences.

**Note 2:** Given `sentiment_analysis` can take a while to run, feel free to create an additional cell when working with the sentiment scores. Once you've come up with your solution, please consolidate your code into one cell and delete the additional cell created to avoid any autograder issues.

```
[76]: sentiment_analysis(msft_news_2010['first_sentence'][0])
```

```
[76]: [{'label': 'NEGATIVE', 'score': 0.9986212253570557}]
```

```
[77]: sentimentos = []
score = []

for i in msft_news_2010['first_sentence'].astype(str):
    sentimentos.append(sentiment_analysis(i[0])[0])

for i in sentimentos:
```



score

17

0.7481208443641663,  
0.9797317385673523,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.985403299331665,  
0.985403299331665,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.5551302433013916,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.9309634566307068,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9845702052116394,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,

0.9475017189979553,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
-0.7622151970863342,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9797317385673523,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
-0.9681645035743713,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.9806599020957947,  
0.9255135655403137,  
0.985403299331665,  
0.985403299331665,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.985403299331665,  
0.985403299331665,

0.985403299331665,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9881240725517273,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
-0.9681645035743713,  
0.9129437208175659,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.9872909188270569,  
0.7481208443641663,  
0.7481208443641663,

0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.985403299331665,  
0.985403299331665,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9918028712272644,  
0.7481208443641663,  
0.7481208443641663,  
0.9309634566307068,  
0.9309634566307068,  
0.7481208443641663,  
0.985403299331665,  
0.9909850358963013,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9894697666168213,  
0.9255135655403137,  
0.7481208443641663,

0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9909850358963013,  
0.9475017189979553,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.9255135655403137,  
0.9872909188270569,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
-0.595452070236206,  
0.7481208443641663,  
0.7481208443641663,  
0.9401155710220337,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9309634566307068,

0.7481208443641663,  
-0.7622151970863342,  
0.7481208443641663,  
0.7481208443641663,  
0.9506378173828125,  
0.7481208443641663,  
0.9475017189979553,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
-0.7622151970863342,  
0.7481208443641663,  
0.8923606872558594,  
0.9475017189979553,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.814979076385498,  
0.985403299331665,  
-0.9681645035743713,  
0.9309634566307068,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9845702052116394,  
0.9775683879852295,  
0.7481208443641663,  
0.9506378173828125,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9845702052116394,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,

0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9845702052116394,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.9309634566307068,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.9475017189979553,



0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.9806599020957947,  
0.7481208443641663,  
0.7481208443641663,  
0.9309634566307068,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9506378173828125,  
0.9918028712272644,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.8923606872558594,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.9255135655403137,

0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9309634566307068,  
0.7481208443641663,  
0.9475017189979553,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9833804368972778,  
0.9475017189979553,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.9129437208175659,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.985403299331665,  
0.8923606872558594,  
0.7481208443641663,  
0.985403299331665,  
0.985403299331665,  
0.7481208443641663,  
0.9861552119255066,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.9918028712272644,  
0.7481208443641663,

0.7481208443641663,  
0.7481208443641663,  
0.985403299331665,  
0.985403299331665,  
0.985403299331665,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.9475017189979553,  
0.9820075631141663,  
0.9861552119255066,  
0.9845702052116394,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
-0.595452070236206,  
0.7481208443641663,  
0.9762058258056641,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,

0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9845702052116394,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.7481208443641663,  
0.9255135655403137,  
0.7481208443641663,  
0.7481208443641663,  
0.9475017189979553,  
0.9894697666168213,  
0.7481208443641663,  
0.9475017189979553,  
0.9872909188270569,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.7481208443641663,  
0.8923606872558594,  
0.7481208443641663,  
0.9762058258056641,  
0.7481208443641663,

```

0.7481208443641663,
0.8923606872558594,
0.7481208443641663,
0.7481208443641663,
0.7481208443641663,
0.985403299331665,
0.985403299331665,
0.985403299331665,
0.8923606872558594,
0.7481208443641663,
0.7481208443641663,
0.7481208443641663,
0.7481208443641663,
0.9475017189979553,
0.7481208443641663,
0.7481208443641663,
0.7481208443641663,
0.7481208443641663,
0.7481208443641663,
0.9475017189979553,
0.7481208443641663]

```

```

[78]: msft_news_2010['sentence_sentiment'] = score
      msft_news_2010.head(1)

```

```

[78]:      id                                     title \
0  95357231  Netflix Profit Jumps 44% on New Users

      released_at \
0  <date>April 21 2010</date> <time>23:52:36</time>

      content \
0  netflix inc. said first-quarter\nprofit rose 44 percent as the movie
subscription service signed\nup new customers and increased online offerings.
\n net income advanced to $32.3 million, or 59 cents a share,\nfrom $22.4
million, or 37 cents, a year earlier, the los gatos,\n...

      path \
0  ./2008_2012_msft_bloomberg_news/netflix-quarterly-profit-increases-44-as-
movie-rental-service-adds-users.txt

      Month Day Year Hour Minute Second \
0  April   21  2010   23     52     36

      first_sentence \
0  \n "if we had offered a pay-per-view service for new\nreleases, we would be
in conflict with a broad range of\ncompanies, including wal-mart, microsoft,

```

sony and apple,"\\nhastings said.

```
    sentence_sentiment
0          0.748121
```

---

### 1.5.2 Question 3b

We can now turn to an alternative, more accurate way of determining the sentiment score of articles — getting the sentiment based on the entire text, rather than getting sentiment based on the first sentence including “microsoft” or “msft” in the text. Let’s load in `data/article_sentiment_logs.csv`, which contains sentiment scores of the full articles as a `DataFrame` `full_sentiments`. In this file, you are provided with logs which include the `id`, `score`, and `label` (“N” for “NEGATIVE” and “P” for “POSITIVE”) in the following format:

```
<device:1> <id:77243971> <result: [0.9963290095329285 (N)]>
<device:0> <id:14799046> <result: [0.9980687499046326 (N)]>
<device:1> <id:43064156> <result: [0.997868537902832 (N)]>
<device:1> <id:29402508> <result: [0.9924335479736328 (N)]>
...
```

Run the following cell to load in the `DataFrame` and see what it contains:

```
[79]: # Run this cell; no further action is needed
full_sentiments = pd.read_csv('/content/article_sentiment_logs.csv')
full_sentiments.head()
```

```
[79]:   RunNum                                     log
0      0  <device:0> <id:77243971> <result: [0.9963290095329285 (N)]>
1      1  <device:0> <id:14799046> <result: [0.9980687499046326 (N)]>
2      2  <device:0> <id:43064156> <result: [0.997868537902832 (N)]>
3      3  <device:0> <id:29402508> <result: [0.9924335479736328 (N)]>
4      4  <device:0> <id:71427879> <result: [0.9897157549858093 (N)]>
```

Using the logs, modify `full_sentiments` so it ultimately just contains the `id` and `content_score` (a number ranging from -1 to 1). Then, merge this with `msft_news_2010` so we can see the results of our two methods of calculating sentiment side by side. Assign this merged `DataFrame` to `msft_scores_2010`. After the merge, make sure that only articles from 2010 appear and that the index of the `DataFrame` is the article `id`.

**Note 1:** You need to negate the score of negatively classified articles (indicated by “N”).

**Note 2:** If you run into issues when merging, you may need to reset `full_sentiments` by running the above cell again.

**Hint 1:** The articles have a primary key `id`.

**Hint 2:** Feel free to reference how you calculated sentiment scores in `q3a1i`.

```
[80]: teste = pd.DataFrame()

teste['id_senti'] = full_sentiments['log'].str.extract(r'<device:0> <id:(\d+)').
↳dropna()
teste['id_senti'] = teste['id_senti'].astype(int)
teste['content_score'] = full_sentiments['log'].str.extract(r'<device:0> <id:
↳\d+> <result: \[(\d+\.\d+)\)').

full_sentiments = teste
full_sentiments
```

```
[80]:
```

	id_senti	content_score
0	77243971	0.9963290095329285
1	14799046	0.9980687499046326
2	43064156	0.997868537902832
3	29402508	0.9924335479736328
4	71427879	0.9897157549858093
...	...	...
4626	44194854	0.9751061201095581
4627	84449274	0.994696855545044
4628	26925649	0.9863374829292297
4632	34512603	0.9963667392730713
4633	39595609	0.7977901101112366

[2320 rows x 2 columns]

```
[81]: msft_scores_2010 = pd.concat([msft_news_2010, full_sentiments], axis=1)
msft_scores_2010 = msft_scores_2010.set_index('id')
msft_scores_2010.head(1)
```

```
[81]:
```

	title \
id	
95357231.0	Netflix Profit Jumps 44% on New Users
	released_at \
id	
95357231.0	<date>April 21 2010</date> <time>23:52:36</time>
	content \
id	
95357231.0	netflix inc. said first-quarter\nprofit rose 44 percent as the movie subscription service signed\nup new customers and increased online offerings. \n net income advanced to \$32.3 million, or 59 cents a share,\nfrom \$22.4 million, or 37 cents, a year earlier, the los gatos,\nn...
	path \
id	

```
95357231.0 ./2008_2012_msft_bloomberg_news/netflix-quarterly-profit-
increases-44-as-movie-rental-service-adds-users.txt
```

```

Month Day Year Hour Minute Second \
id
95357231.0 April 21.0 2010.0 23.0 52.0 36.0
```

```

first_sentence \
id
95357231.0 \n "if we had offered a pay-per-view service for new\nreleases, we
would be in conflict with a broad range of\ncompanies, including wal-mart,
microsoft, sony and apple," \nhastings said.
```

```

sentence_sentiment id_senti content_score
id
95357231.0 0.748121 77243971.0 0.9963290095329285
```

### 1.5.3 Question 3c

Let's dive deeper into our two methods of calculating sentiment and analyze the accuracy of the method used in q3b.

**Question 3c, Part i** Calculate the difference between `content_score` and `sentence_sentiment`. Create a new column `sentiment_difference` in our DataFrame `msft_scores_2010` with these values.

```
[82]: msft_scores_2010['sentiment_difference'] = msft_scores_2010['content_score'].
      ↪ astype(float) - msft_scores_2010['sentence_sentiment']
msft_scores_2010.head(1)
```

```
[82]:
title \
id
95357231.0 Netflix Profit Jumps 44% on New Users

released_at \
id
95357231.0 <date>April 21 2010</date> <time>23:52:36</time>

content \
id
95357231.0 netflix inc. said first-quarter\nprofit rose 44 percent as the movie
subscription service signed\nup new customers and increased online offerings.
\n net income advanced to $32.3 million, or 59 cents a share,\nfrom $22.4
million, or 37 cents, a year earlier, the los gatos,\n...
```

```
path \
```



```
id
95357231.0  ./2008_2012_msft_bloomberg_news/netflix-quarterly-profit-
increases-44-as-movie-rental-service-adds-users.txt
```

```
Month Day Year Hour Minute Second \
id
95357231.0 April 21.0 2010.0 23.0 52.0 36.0
```

```
first_sentence \
id
95357231.0  \n "if we had offered a pay-per-view service for new\nreleases, we
would be in conflict with a broad range of\ncompanies, including wal-mart,
microsoft, sony and apple,"\nhastings said.
```


```
sentence_sentiment id_senti content_score \
id
95357231.0 0.748121 77243971.0 0.9963290095329285
```

```
sentiment_difference
id
95357231.0 0.248208
```

**Question 3c, Part ii** Below we have provided a plot looking at these differences. Comment on why we see differences when calculating the sentiment of an article as the sentiment of the first sentence mentioning “microsoft” or “msft” in the article versus the sentiment of the entire article itself. How does context play a role when evaluating the sentiment of a text?

```
[85]: sns.kdeplot(msft_scores_2010['sentiment_difference'])
plt.xlabel('Sentiment difference')
plt.title('Difference between full and approximate sentiment scores');
```

Difference between full and approximate sentiment scores



A diferença ocorre porque uma única frase não define um texto inteiro. Um texto deve sempre ser lido do começo ao fim para que garantir que a uma frase não esteja fora de contexto.

A célula abaixo não funciona, ela está em modo de leitura e não consigo nem modificar e nem apagar

```
[84]: # Run this cell; no further action is needed
full_sentiments = pd.read_csv('data/article_sentiment_logs.csv')
full_sentiments.head()
```

```
-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-84-b65f36a4dd87> in <cell line: 0>()
      1 # Run this cell; no further action is needed
----> 2 full_sentiments = pd.read_csv('data/article_sentiment_logs.csv')
      3 full_sentiments.head()

/usr/local/lib/python3.11/dist-packages/pandas/io/parsers/readers.py in
    ↳ read_csv(filepath_or_buffer, sep, delimiter, header, names, index_col,
    ↳ usecols, dtype, engine, converters, true_values, false_values,
    ↳ skipinitialspace, skiprows, skipfooter, nrows, na_values, keep_default_na,
    ↳ na_filter, verbose, skip_blank_lines, parse_dates, infer_datetime_format,
    ↳ keep_date_col, date_parser, date_format, dayfirst, cache_dates, iterator,
    ↳ chunksize, compression, thousands, decimal, lineterminator, quotechar,
    ↳ quoting, doublequote, escapechar, comment, encoding, encoding_errors, dialect,
    ↳ on_bad_lines, delim_whitespace, low_memory, memory_map, float_precision,
    ↳ storage_options, dtype_backend)
```

```

1024     kwds.update(kwds_defaults)
1025
-> 1026     return _read(filepath_or_buffer, kwds)
1027
1028

/usr/local/lib/python3.11/dist-packages/pandas/io/parsers/readers.py in
↳ _read(filepath_or_buffer, kwds)
    618
    619     # Create the parser.
--> 620     parser = TextFileReader(filepath_or_buffer, **kwds)
    621
    622     if chunksize or iterator:

/usr/local/lib/python3.11/dist-packages/pandas/io/parsers/readers.py in
↳ __init__(self, f, engine, **kwds)
    1618
    1619         self.handles: IOHandles | None = None
-> 1620         self._engine = self._make_engine(f, self.engine)
    1621
    1622         def close(self) -> None:

/usr/local/lib/python3.11/dist-packages/pandas/io/parsers/readers.py in
↳ _make_engine(self, f, engine)
    1878             if "b" not in mode:
    1879                 mode += "b"
-> 1880                 self.handles = get_handle(

    1881                     f,
    1882                     mode,

/usr/local/lib/python3.11/dist-packages/pandas/io/common.py in
↳ get_handle(path_or_buf, mode, encoding, compression, memory_map, is_text,
↳ errors, storage_options)
    871         if ioargs.encoding and "b" not in ioargs.mode:
    872             # Encoding
--> 873             handle = open(

    874                 handle,
    875                 ioargs.mode,

FileNotFoundError: [Errno 2] No such file or directory: 'data/
↳ article_sentiment_logs.csv'

```